

# TiPAI: Tournament Inpainting for Patch-Level Alignment in Text-to-Image

Anonymous ACL submission

## Abstract

We propose **TiPAI-TSPO**, a decoding-time framework that makes text-to-image generation both *faithful* and *policy-compliant* by editing only where it matters. Starting from a base diffuser, each timestep exposes a low-cost audit view; we mine a small set of suspect regions and run a *tournament of local inpaints* ( $N = 5$ ) produced by an auditor–inpaint model. A learned *auditor–scorer*—trained on 100k DETONATE “selected vs. rejected” pairs—assigns a composite score that integrates text–image faithfulness, policy safety, and seam quality. A *control guard* accepts an edit only if it *beats the unedited control by a calibrated margin* and passes per-class safety/faithfulness thresholds, yielding *monotone non-regression* at the patch level.

To reduce retries and latency, we introduce *Tournament Sampling Policy Optimization (TSPO)*: a light policy over generator knobs (e.g., mask dilation, CFG-inside-mask, latent noise jitter, short-inversion depth) trained with listwise credits and diversity/compute regularizers. TSPO learns to propose winners more often with less compute. We couple edits to the diffusion trajectory via timestep-aware inpainting, short DDIM inversion, and latent-space blending to avoid visual scars. A final calibration stage turns raw scores into reliable acceptance decisions with interpretable knobs (margin  $\delta$ , policy  $\tau_P$ , faithfulness  $\tau_F$ ) scheduled across timesteps.

TiPAI-TSPO provides a practical, plug-and-play route to *aligned, efficient* text-to-image decoding.

## 1 Technical Approach (Step by Step)

We align text-to-image (T2I) decoding by performing *local, timestep-aware edits* that are accepted only when they strictly improve faithfulness and policy safety over the unedited control. The system

has three training stages: (A) an auditor–scorer  $A_s$  trained on DETONATE pairs to produce global/patch scores and risk maps; (B) an auditor–inpaint editor  $A_g$  pretrained to rewrite local regions, then coupled with a best-of- $N$  tournament and a learned sampling policy (TSPO); and (C) a calibration stage that turns raw scores into reliable decisions with interpretable thresholds.

### Preliminaries and Notation

- Prompt  $p$ , diffusion trajectory  $\{x_t\}_{t=0}^T$  or latent  $\{z_t\}$  (VAE space).
- Base model produces control proposal at step  $t$ :  $x_{t-1}^{\text{ctrl}} = \Phi_{\text{base}}(x_t, p, t)$  (or  $z_{t-1}^{\text{ctrl}}$ ).
- An *audit view*  $I_{t-1}$  is decoded (possibly at a lower resolution) for patch mining and scoring.
- From  $I_{t-1}$ , mine  $K_t$  regions  $\mathcal{R}_t = \{(R_k, m_k)\}$  using fused saliency: CLIP Grad-CAM drift, detector priors (NSFW/weapon/symbol/logo/age), OCR, and feature-difference heatmaps. Each region has a binary (feathered) mask  $m$ ; crops are  $I_{t-1,R}$ .

### Stage A: Auditor–Scorer $A_s$ (Global + Patch Ranking with Risk Heads)

**Goal.** Learn scalar scores  $S(p, I)$  and  $S_R(p, I_R)$  that order human-*selected* images above *rejected* ones, and produce per-class risk maps for policy guards.

### Encoders and Heads.

$$S(p, I) = h_\theta(E_t(p), E_i(I)), \quad S_R(p, I_R) = h_\theta^{\text{patch}}(E_t(p), E_i(I_R)),$$

with a lightweight decoder that predicts per-class heatmaps  $r_c \in [0, 1]^{H \times W}$  from shared features.

**Training Data.** Pairs  $\{(p, I_+, I_-)\}$  from DETONATE. For each pair, mine  $K$  patches  $R_k$  and assemble  $(p, I_{\pm, R_k}, m_{R_k})$ . Optional reason codes  $y_c \in \{0, 1\}$  per policy class.

## Losses (no numbering).

$$\mathcal{L}_{\text{pair}} = \frac{1}{N} \sum_n \log(1 + \exp(-(S_+^{(n)} - S_-^{(n)}))),$$

$$\mathcal{L}_{\text{patch}} = \frac{1}{N} \sum_n \frac{1}{K_n} \sum_k \log(1 + \exp(-(S_{R_k,+}^{(n)} - S_{R_k,-}^{(n)}))),$$

$$\mathcal{L}_{\text{policy}} = \frac{1}{N} \sum_n \sum_c \text{BCE}\left(\sigma(\bar{\ell}_c(I^{(n)})), y_c^{(n)}\right), \quad \mathcal{L}_{\text{sal}} = \frac{1}{N} \sum_n \sum_c \left(1 - \text{Dice}(r_c^{(n)}, \bar{m}^{(n)})\right),$$

$$\mathcal{L}_A = \lambda_1 \mathcal{L}_{\text{pair}} + \lambda_2 \mathcal{L}_{\text{patch}} + \lambda_3 \mathcal{L}_{\text{policy}} + \lambda_4 \mathcal{L}_{\text{sal}}.$$

## Step-by-step recipe.

1. Encode  $(p, I)$  and  $(p, I_R)$  to obtain global and patch features.
2. Predict  $S, S_R$  and risk maps  $r_c$ ; train with  $\mathcal{L}_A$  using a curriculum (start with high-margin pairs, then anneal).
3. Validate with Global Pair-AUC, Patch Pair-AUC, per-class ROC-AUC, and ECE for  $S$ .

*Outputs used later:* calibrated-ready  $S, S_R$ , risk maps  $r_c$ , reason vector.

## Stage B: Auditor–Inpaint $A_g$ + Tournament Sampling Policy Optimization (TSPO)

### B1. Pretrain $A_g$ as a timestep-aware local editor. Input formation (context + noise inside mask).

$$X = I_{t-1} \odot (1 - m) + \epsilon \odot m, \quad \epsilon \sim \mathcal{N}(0, \sigma_t^2), \quad t \in \text{chosen timestep band}.$$

**Target.**  $Y = I_R^+$  from the DETONATE selected image.

**Inpaint prediction.**  $\hat{I}_R = A_g(X, m, p, t)$  in RGB or latent.

### Losses.

$$\mathcal{L}_{\text{inpaint}} = \beta_1 \|\hat{I}_R - Y\|_1 + \beta_2 (1 - \text{CLIP}(p, \hat{I}_R)) + \beta_3 \text{Risk}(\hat{I}_R) + \beta_4 \text{LPIPS}_{\partial m}(\hat{I}_R, Y),$$

$$\mathcal{L}_{\text{pref}} = -\log \frac{\exp(\langle f(\hat{I}_R), f(I_R^+) \rangle / \tau)}{\exp(\langle f(\hat{I}_R), f(I_R^-) \rangle / \tau) + \exp(\langle f(\hat{I}_R), f(I_R) \rangle / \tau)}, \quad \mathcal{L}_G^{\text{pre}} = \mathcal{L}_{\text{inpaint}} + \beta_5 \mathcal{L}_{\text{pref}}.$$

**Short inversion and latent blending (seam-free coupling).** If editing in RGB, map  $\hat{I}_R$  back to the scheduler state using a short DDIM inversion of  $d$  steps, then blend in latent:

$$z_{t-1} \leftarrow (1 - \alpha m) \odot z_{t-1}^{\text{ctrl}} + \alpha m \odot z_{t-1,R}^{\text{edit}}, \quad \alpha = \alpha(t) \in [0, 1],$$

with feathered  $m$  and timestep-aware  $\alpha$ .

### Pretrain steps.

1. Sample a pair  $(p, I^+, I^-)$ , mine  $R, m$ , choose a timestep  $t$  and form  $X$ .  
112  
113
  2. Predict  $\hat{I}_R = A_g(X, m, p, t)$ ; compute  $\mathcal{L}_G^{\text{pre}}$ ; update  $A_g$ .  
114  
115
  3. Validate with patch risk reduction  $\Delta r$ ,  $\Delta \text{CLIP}$  on crops, boundary LPIPS on a ring.  
116  
117
- B2. Define the tournament and guards (best-of-N + control). Candidate generation.** For each region, draw  $N$  actions  $a_i$  from a generator policy  $\pi_\theta(a | s)$  with state  $s = (p, z_{t-1}, I_{t-1}, m, t)$ ; produce candidates  $C_i = A_g(s; a_i)$ . Always include the unedited control  $C_0$ .  
118  
119
- Scoring.** Use  $A_s$  to produce, for each  $i$ ,  
120  
121  
122  
123
- $$(S_i, F_i, P_i, B_i) = A_s(p, \text{Compose}(I_{t-1}, C_i, R)),$$
- where  $F$  is faithfulness,  $P$  is policy safety (risk complement),  $B$  is seam quality (e.g.,  $B = \exp(-\kappa \text{LPIPS}_{\partial m})$ ).  
124  
125  
126  
127  
128
- Guarded margin utility and selection.**  
129
- $$u_i = (S_i - S_0 - \delta)_+ \cdot \mathbf{1}[P_i \geq \tau_P(t)] \cdot \mathbf{1}[F_i \geq \tau_F(t)] \cdot B_i.$$
- Let  $i = \arg \max_i u_i$ . Accept  $C_i$  iff  $u_i > 0$ ; otherwise keep  $C_0$  (*control guard* ensures non-regression).  
130  
131  
132  
133
- B3. Learn the sampling policy: TSPO. Actions (knobs).** Mask dilation/feathering, inside-mask CFG, prompt token emphasis, latent noise jitter, dropout/seed, short inversion depth  $d$ , light LoRA routing.  
134  
135  
136  
137  
138
- Credits.** Leave-one-out advantage for each candidate:  
139  
140
- $$A_i = u_i - \max_{j \neq i} u_j, \quad \text{or soft credits } w_i = \text{softmax}(u_i / \tau) - \frac{1}{N}.$$
- Objective (policy gradient, no numbering).**  
141  
142
- $$\mathcal{L}_{\text{TSPO}} = -\sum_{i=1}^N w_i \log \pi_\theta(a_i | s) - \beta H[\pi_\theta(\cdot | s)] + \lambda_c \text{Cost}(\{a_i\}) - \lambda_{\text{div}} \sum_{1 \leq i < j \leq N} d(C_i, C_j),$$
- where **Cost** penalizes short inversions/extraneous decodes;  $d$  is a masked feature distance encouraging within-tournament diversity.  
143  
144  
145  
146
- TSPO steps.**  
147
1. For each audited region, sample  $N$  actions, generate  $C_{1:N}$ , compute  $u_{0:N}$  and pick  $i$ .  
148  
149
  2. Accumulate  $(a_{1:N}, w_{1:N}, \text{Cost}, d)$  and update  $\theta$  by minimizing  $\mathcal{L}_{\text{TSPO}}$  (REINFORCE with entropy/compute/diversity regularizers).  
150  
151  
152
  3. (Optional) Improve the judge listwise with a Plackett–Luce/ListNet loss over  $\hat{S}_i$  against a target distribution concentrated on the winner.  
153  
154  
155

156 **Stage C: Calibration  $\Rightarrow$  Reliable Decisions**

157 **Score calibration on held-out tournaments.** Fit  
 158 isotonic or Platt mapping from raw  $S$  to probability  
 159  $\hat{p}$  on held-out data:

160 
$$\hat{p} = \sigma\left(\frac{S - b}{T}\right) \quad \text{or} \quad \hat{p} = \text{Iso}(S),$$

161 choosing parameters to minimize NLL and ECE.

162 **Choose operating points (interpretable knobs).**

- 163 • Margin  $\delta$ : the smallest  $\Delta$  such that  $\Pr[\text{true win} | S_{\text{cand}} \geq S_{\text{ctrl}} + \Delta] \geq 0.9$ .
- 164
- 165 • Policy thresholds  $\tau_P(t)$ : per-class ROC operating  
 166 points (stricter late).
- 167
- 168 • Faithfulness  $\tau_F(t)$ : stricter mid-trajectory;  
 169 slightly relaxed at the very end to avoid over-  
 170 sanitization.
- 171
- Seam weight  $\kappa$  and artifact budget (max boundary  
 LPIPS).

172 **Deployed decision rule (per patch, per step).**

- 173 1. Sample  $N$  candidates via  $\pi_\theta$ ; score with calibrated  
 174  $A_s$  to obtain  $(S_i, F_i, P_i, B_i)$ .
- 175 2. If there exists a candidate with  $S_i \geq S_0 + \delta$ ,  
 176  $P_i \geq \tau_P(t)$ ,  $F_i \geq \tau_F(t)$ , and  $B_i$  above the seam  
 177 threshold, accept the best one; else keep control  
 178 and optionally re-audit earlier.
- 179 3. Log tournaments for continual TSPO updates and  
 180 periodic recalibration.

181 **Coupling to the Diffusion Scheduler**  
 182 **(Seam-Free Edits)**

- 183 1. Run audits on a subset  $\mathcal{T}_{\text{audit}}$  of timesteps (e.g., ev-  
 184 ery third step), coarse scales early and fine scales  
 185 late.
- 186 2. If  $A_g$  edits in RGB, use a short deterministic  
 187 DDIM inversion of  $d$  steps to re-land  $\hat{I}_R$  into  
 188  $z_{t-1,R}^{\text{edit}}$  consistent with the scheduler's  $(\bar{\alpha}_t)$  coeffi-  
 189 cients.
- 190 3. Blend in latent with feathered masks and timestep-  
 191 aware  $\alpha(t)$ ; decode only once per step for the audit  
 192 view; reuse feature pyramids across candidates.

Complexity and Compute–Quality Tradeoff

For each audited step: cost scales as

$$\mathcal{O}(K_t \cdot (N \cdot C_{A_g} + C_{A_s} + d \cdot C_{\text{DDIMInv}})).$$

Batch inpainting and scoring; cache text/image fea-  
 tures. TSPO learns to reduce expected inversions  
 and favors low-cost actions unless higher-cost ac-  
 tions increase the guarded utility.

**Properties (Intuition Sketches, No Numbers)**

- **Monotone non-regression (patchwise).** The control guard accepts only strict improvements over control that satisfy policy/faithfulness thresholds; otherwise it returns the control, so guarded objectives do not decrease locally.
- **TSPO convergence (fixed judge).** With bounded variance and entropy regularization, policy gradients over a finite action space converge to a stationary point of the expected listwise utility; the compute penalty forms a Lagrangian that places the learned policy on a quality–latency Pareto frontier.
- **Calibration reliability.** Isotonic/Platt calibration on held-out tournaments reduces ECE so fixed  $(\delta, \tau_P, \tau_F)$  maintain target operating characteristics across prompts/seeds drawn from the same regime.

**Algorithm (Annotated, Step by Step)**

1. **Control step:** produce  $z_{t-1}^{\text{ctrl}}$  (or  $x_{t-1}^{\text{ctrl}}$ ) and decode  $I_{t-1}$ .
2. **Mine regions:** build  $\mathcal{R}_t = \{(R_k, m_k)\}$  using fused saliency signals.
3. **For each region**  $(R, m)$ :
  - (a) **Generate candidates:** sample  $N$  actions  $a_i \sim \pi_\theta(\cdot | s)$ ; produce  $C_{1:N}$  with  $A_g$ ; add  $C_0$  (control).
  - (b) **Score:** compute  $(S_i, F_i, P_i, B_i)$  with  $A_s$  (batched).
  - (c) **Select:** compute  $u_i$  with the guarded margin; if a candidate wins, short-invert and latent-blend into  $z_{t-1}$ ; else keep control.
  - (d) **Log:** store  $(a_{1:N}, u_{0:N}, \text{cost}, \text{winner})$  for TSPO.
4. **TSPO update (online or mini-batch):** minimize  $\mathcal{L}_{\text{TSPO}}$  with entropy/compute/diversity terms; periodically refine the judge listwise and re-calibrate scores on a held-out slice.

**References**